

Date - 03.11.23

* Producer Consumer problem :-

int count = 0; ① Producer
{ void producer (void)
{

int item;
while (true)

{
 producer-item (item);
 while (count == N); // Buffer full
 Buffer[in] = item;
 in = in + 1 mod N
 count = count ++;
}

void consumer (void) { ② Consumer

int item c;
while (true)

{ while (count == 0); // Buffer empty


item c = Buffer[out];

out = out + 1 mod N;

count = count --;

Process Item (Item c);

}



* count++ :-

- ① Load R_p $m(\text{count})$
- ② INCR R_p 4 *
- ③ $m(\text{count}), R_p$

* count-- :-

- ① Load R_c $m(\text{count})$
- ② DECR R_c
- ③ $m(\text{count}), R_c$

→ The count is a shared variable that tells us the total no. of items present in the buffer. Item 'p' is variable that denotes items to produced, 'N' is the total buffer size. 'in' is the memory location where item to be produced and 'out' is the location from where item is to be consumed.

* Counting Semaphore

① Producer

Binary Semaphore $S = 1$

Produce item (item p)

Down (empty) — x (context switch happens)

Down (S);

Buffer[in] = item p; CS

in = in + 1 mod N

up(S)

up(full)

empty = 5/5

full = 3/3

② Consumer

Down(full);

Down(S);

Item C = Buffer[out]

out = out + 1 mod N


up(S)

up(empty) CS

ME ✓

~~date~~ Date - 7/11/23

* Reader - Writer Problem :-



Database

② B Semaphore Mutex = 0, db = 1;

```
int rc = 0;
void writer(void) {
    while(True);
    Down(db);
    <Access Database>
    Up(db);
}
```

```
void Reader(void)
```

```
{ while(true);
```

```
  { down(mutex)
```

```
    rc = rc + 1
```

```
    if (rc == 1)
```

```
      { down(db)
```

```
        up(mutex)
```

```
        <Access database>
```

```
        down(mutex)
```

```
        rc = rc - 1
```

→ The problem in Reader - Writer's -

There is a data area shared amongst no. of processes. The data area could be a block, a file or a bank of processor register. There are no. of processes that only read the data area (reader's) and a no. of processes that only write to the data area (writer's). The condⁿ that must be satisfied is as follows -

① Any no. of readers may simultaneously read the file.

② Only one writer at a time may write to the file.

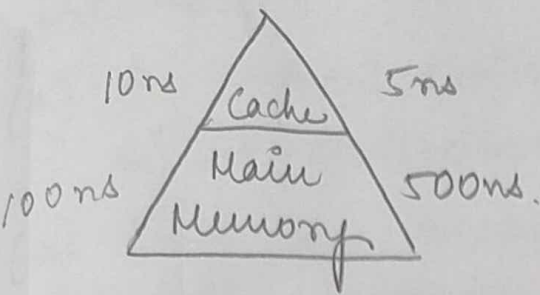
③ If the writer is writing to the file, no reader may read it. Therefore, reader and writer processes that are required to exclude one another and writer are processes that are required to ~~all~~ exclude all other processes, i.e. both readers & writers.

* The solution is provided using semaphores showing one instance each of a reader or a writer. The solution doesn't change for multiple readers & writer. The writer process is simple.

The semaphore db is used to ensure mutual exclusion. As long as one writer is accessing the shared data area, no other reader & writer may access it. The reader process also uses mutex & db to ensure mutual exclusion. However to allow multiple readers, we require that, when there are no readers are reading, the first reader that attempts to read should wait on mutex. When ~~there is~~ ^{there is} at least one reader reading, subsequent readers need not wait before entering. The global variable read count rc is used to keep track of the no. of readers and semaphore is used to ensure that read count is updated properly.

Hit ratio - 0.95
Miss ratio - 0.05

$$\begin{aligned} \text{Access time} &= 0.95 \times 10 \text{ ns} + 0.05 \times 100 \text{ ns} \\ &= 9.50 \text{ ns} + 5.00 \text{ ns} \\ &= 9.5 + 5.0 \\ &= 14.5 \text{ ns} \end{aligned}$$



Hit ratio = 0.8

$$\begin{aligned} &= 0.8 \times 5 + 0.2 \times 500 \\ &= 4.0 + 100.0 \\ &= 104 \text{ ns (Total access time)} \end{aligned}$$

Cache = 30 ns
MM = 400 ns

Hit ratio = 0.7

$$\begin{aligned} \text{TAT} &= 30 \times 0.7 + 0.3 \times 400 = 21.0 + 120.0 \\ &= 120 + 21 = 141 \text{ ns} \end{aligned}$$

80

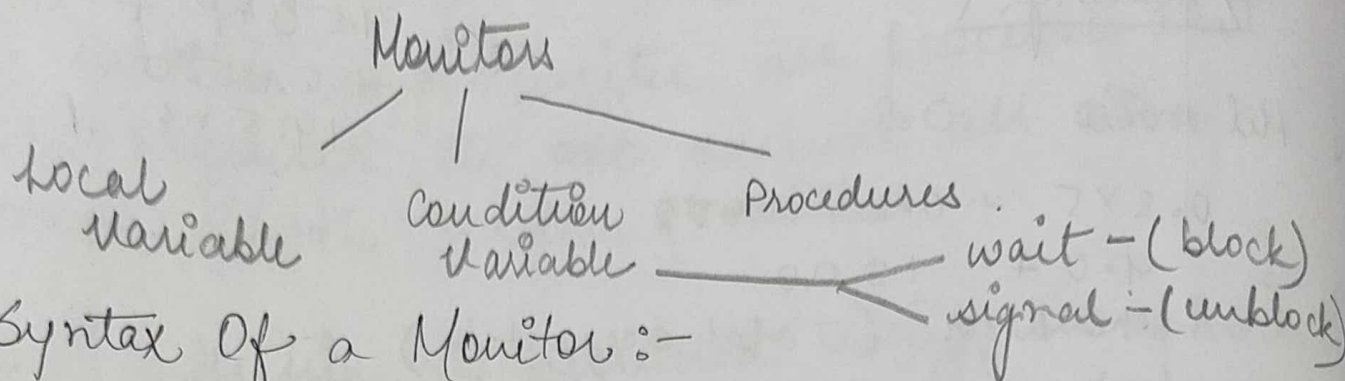
$$\begin{aligned} \text{Total time} &= 0.8 (\text{time for read}) + 0.2 (\text{time for write}) \\ &= 0.8 \times 30 \times 0.7 + (0.2 \times 0.3 \times 400 + 0.2 \times 0.7 \times 30) \\ &= 16.80 + 24.00 + 4.2 \\ &= 16.8 + 24.0 + 4.2 \\ &= 40.8 \text{ ns} + 4.2 \\ &= 45.0 \text{ ns} \end{aligned}$$

$$\begin{array}{r} 16 \\ + 3 \\ \hline 19 \\ 168 \\ \hline 14 \\ + 12 \\ \hline 420 \end{array}$$

The local data variables are accessible only by the procedure & not by external reference & local data. The characteristic of one or more procedures that are organized as a separate module consisting of a procedure & its associated data is called a subprogram. A subprogram is a module that can be called from other modules. It is a self-contained unit that can be used repeatedly in a program. It is a way of organizing a program into smaller, more manageable parts. It is a way of organizing a program into smaller, more manageable parts. It is a way of organizing a program into smaller, more manageable parts.

Date - 8.11.23

* Monitors :-



* Syntax Of a Monitor :-

monitor monitor-name

{

// local variables

// condition variables

Procedure P₁()

{
=
}

Procedure P₂()

{

Procedure P_n()

{
=
}

}

// Initialization Code

}

→ A monitor is a software module consisting of one or more procedures, an initialization sequence & local data. The characteristics are-

(1) The local data variables are accessible only by monitor procedure & not by external

- no count
- (2) A process enters the monitor by invoking one of its procedure.
 - (3) Only one process can execute monitor at a time. Any other processes that have invoked the procedure, monitor have blocked the monitor to be available.

Monitors support synchronization by the use of condⁿ variables that are present in the monitor and are accessible only within the monitor.

The 2 functions of condⁿ variables are wait & signal. The operation wait means that the process invoking this operation is suspended until another process invokes signal. The signal operation assumes exactly one suspended process.

*:

monitor producer-consumer

{
 condition variables full, empty;
 int count = 0;
}

Procedure enter-item()

{
 if (count == N) {
 wait(full);
 }
 enter-item(item P);
 count = count + 1;
 if (count == 1) {
 signal(empty);
 }
}

Procedure remove-item () {

if (count == 0) {

wait(empty);

remove(item);

count = count - 1;

if (count == N-1) {

signal(full);

}

}

Procedure Producer ()

{

while (true)

produce item (item p)

producer-consumer.enter-item;

}

Procedure Consumer ()

{

while (True)

producer-consumer.remove-item;

consume item (item c);


}

Date





→ In context of producer-consumer problem, it is not possible for both producer and consumer to simultaneously access the buffer. The programmer places appropriate wait & signal ^{primitives} inside the monitor to prevent processes from depositing items in a full buffer and removing them from an empty one. Here we have used full & empty as condⁿ variables on which wait & signal operⁿ are performed. If the buffer is full, the producer is blocked by issuing wait of full. It is unblocked by consumer by issuing signal of full. If the consumer is blocked by issuing by wait of empty then it is unblocked by producer by issuing signal of empty.

Date - 14.11.23

PNo	AT	BT	CT	TAT	WT=RT
1	1	7	8	7	
2	2	5	16	14	
3	3	1	9	6	
4	4	2	11	97	
5	5	8	24	19	



$$\text{Avg TAT} = 53/5$$



Memory size = 1024 MB

@ Page size = 512 KB

$$\text{No. of pages} = \frac{1024 \times 2^{20}}{512 \times 2^{10}} = \frac{2^{30}}{2^{19}}$$

Page size = Frame size

$$= 2^{11} = 2048 \text{ pages}$$

Data = 16 MB

PS = 4 MB

$$= \frac{2^4 \times 2^{30}}{2^2 \times 2^{20}}$$

$$= \frac{2^2 \times 2^{10}}{2^{12}}$$

* Date - 15.11.23

* Message Passing :-

receive(source, msg)
send(destination, msg)

- (1) Blocking send, Blocking receive
- (2) Unblocking send, Blocking receive
- (3) Unblocking send, Unblocking receive

→ ^{Sender:-} redundancy of message

→ The function of message passing is provided in a form of pair of primitives:-

- (1) Send
- (2) Receive

A process sends information in the form of a message designated by a destination. A process receives information by using receive primitive indicating the source and the message. To achieve synchronization, the receiver can't receive the message until it has been sent by another process. We also need to specify what happens to a process after it issues a send/receive primitive.

The following combinations are possible:-

- (1) Blocking send and blocking receive:-

Both the sender and receiver are blocked until the message is delivered. This combination

allows for strict synchronization.

(2) Nonblocking send and blocking receive:-

The sender may continue to send messages to different processes but the receiver is blocked until the requested message arrives. It helps a process to send one or more messages to a variety of destination as quickly as possible.


(3) Non-blocking send and non-blocking receive:-
Neither sender or receiver is required to wait.

Addressing


Direct

Indirect


1) One-to-one



3) Many-to-one



2) One to many



There are 2 types of addressing :-

1) With direct addressing, the send primitive includes a specific identifier of the destination process. But in case of receive primitive a process must know ahead of time from which process the message is expected. In this case the source parameter of the receive primitive has a value returned, when the receive operation has been performed.

2) Here messages are not sent directly from sender to receiver but rather sent to a shared data structure that consists of queues, that can temporarily hold messages. Such queues are known as mailbox.

The relationship can be :-

1) One-to-one

It allows a private communication link to be setup b/w processes.

2) Many-to-one

It is useful for client server interaction where one process provides services to multiple processes.

3) One-to-many

This allows for one sender and multiple receivers mostly useful in broadcasting of messages.

4) Many-to-many:-

This allows multiple server processes to provide concurrent services to multiple clients.


Note:

In many to one, mailbox is often referred to as port.

* Message Format:-

Message Type	Header
Source ID	
Destination ID	
Message length	
Control Info	Body
Message Content	

* Message passing:-



The Solution we use: if there is a message.

The message is divided into 2 parts :-

① Header - It contains information about the message such as identification of the source and intended destⁿ of the message, The length field and message type to differentiate between different type of messages.

There is control information such as a pointer field ~~which~~ so that we can create a linked list of message.

② Body - The body of the message contains the actual content of the message.

* To ensure mutual exclusion in critical section, we use blocking receive & unblocking send. Different processes share a mailbox which is initialized to contain a single message with all ctrl. content

→ A process wishing to enter a CS, first attempts to receive a message

→ If the ~~message~~ mailbox is empty, the process is blocked.

→ If a process has acquired ^{the} message, it performs work in the CS and place the message back in the mailbox. The message here functions as a token that is passed from process to process.

→ The solnⁿ assumes if there is a message,
it is delivered to only one process and
the others are blocked.

→ If the queue is empty, all the processes
are blocked, until a message is made
available.